


Single-node Hadoop Setup

After installation, we successfully access the port 50070 of localhost.

```
00000000000000000000 of size 328 bytes saved in 0 seconds .
20/09/12 07:07:29 INFO namenode.NameNode.StorageRetentionManager: Going to retain 1 images with txid >= 0
20/09/12 07:07:29 INFO namenode.NameNode: SHUTDOWN_MSG:
/************************************************************
SHUTDOWN_MSG: Shutting down NameNode at instance-1.asia-east2-a.c.driven-plexus-289109.internal/10.170.0.2
************************************************************/
[jessechan@instance-1 hadoop]$ sh bin/start-dfs.sh
Starting namenodes on [localhost]
localhost: starting namenode, logging to /usr/local/hadoop/logs/hadoop-jessechan-namenode-instance-1.out
localhost: datanode running as process 2322. Stop it first.
Starting secondarynamenodes [0.0.0.0]
0.0.0.0: secondarynamenode running as process 2520. Stop it first.
[jessechan@instance-1 hadoop]$ curl http://localhost:50070
<!--
Licensed to the Apache Software Foundation (ASF) under one or more
contributor license agreements.  See the NOTICE file distributed with
this work for additional information regarding copyright ownership.
The ASF licenses this file to You under the Apache license, Version 2.0
(the "License"); you may not use this file except in compliance with
the License.  You may obtain a copy of the License at

    http://www.apache.org/licenses/LICENSE-2.0

Unless required by applicable law or agreed to in writing, software
distributed under the License is distributed on an "AS IS" BASIS,
WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
See the License for the specific language governing permissions and
limitations under the License.
-->
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html xmlns="http://www.w3.org/1999/xhtml">
<head>
<meta http-equiv="REFRESH" content="0;url=dfshealth.html" />
<title>Hadoop Administration</title>
</head>
</html>
[jessechan@instance-1 hadoop]$
```



Overview 'localhost:9000' (active)

Started:	Thu Sep 10 21:27:14 +0800 2020
Version:	2.9.2, r626afbaac31ca687bc28471dc841b06ed2cd704
Compiled:	Tue Nov 13 20:42:00 +0800 2018 by ajsaka from branch-2.9.2
Cluster ID:	CID-d91120c-a1b6-4786-8a26-8a7891b7b743
Block Pool ID:	BP-697813659-10.170.0.2-1989744057794

Summary

Security is off.	
SafeMode is off.	
1 files and directories, 0 blocks = 1 total filesystem object(s).	
Heap Memory used 62.21 MB of 176 MB Heap Memory. Max Heap Memory is 889 MB.	
Non Heap Memory used 38.26 MB of 39.31 MB Committed Non Heap Memory. Max Non Heap Memory is 'unbounded'.	
Configured Capacity:	96.75 GB
DFS Used:	24 KB (0%)
Non DFS Used:	2.73 GB
DFS Remaining:	94 GB (97.16%)
Block Pool Used:	24 KB (0%)
DataNodes usages% (Min/Median/Max/stdDev):	0.00% / 0.00% / 0.00% / 0.00%
Live Nodes	1 (Decommissioned: 0, In Maintenance: 0)
Dead Nodes	0 (Decommissioned: 0, In Maintenance: 0)
Decommissioning Nodes	0
Entering Maintenance Nodes	0
Total Datanode Volume Failures	0 (0 B)
Number of Under-Replicated Blocks	0
Number of Blocks Pending Deletion	0
Block Deletion Start Time	Thu Sep 10 21:27:14 +0800 2020
Last Checkpoint Time	Thu Sep 10 21:20:57 +0800 2020

NameNode Journal Status

Current transaction ID: 1	
Journal Manager	State
FileJournalManager[root/tmp/hadoop-jessechan/dfs/name]	EditLogOutputStream[tmp/hadoop-jessechan/dfs/name/current/edits_inprogress_00000000000000000001]

NameNode Storage

Storage Directory	Type	State
/tmp/hadoop-jessechan/dfs/name	IMAGE_AND_EDITS	Active

DFS Storage Types

Storage Type	Configured Capacity	Capacity Used	Capacity Remaining	Block Pool Used	Nodes in Service
DISK	96.75 GB	24 KB (0%)	96.75 GB (99.99%)	24 KB	1

Hadoop, 2018.

We try to run the Terasort. First, we generate a list of random numbers.


```
3492.40.164 - PuTTY
Total time spent by all reduces in occupied slots (ms)=4791
Total time spent by all map tasks (ms)=15092
Total time spent by all reduce tasks (ms)=4791
Total vcore-milliseconds taken by all map tasks=15092
Total vcore-milliseconds taken by all reduce tasks=4791
Total megabyte-milliseconds taken by all map tasks=15454208
Total megabyte-milliseconds taken by all reduce tasks=4905984

Map-Reduce Framework
  Map input records=100000
  Map output records=100000
  Map output bytes=10200000
  Map output materialized bytes=10400012
  Input split bytes=230
  Combine input records=0
  Combine output records=0
  Reduce input groups=100000
  Reduce shuffle bytes=10400012
  Reduce input records=100000
  Reduce output records=100000
  Spilled Records=200000
  Shuffled Maps=2
  Failed Shuffles=0
  Merged Map outputs=2
  GC time elapsed (ms)=456
  CPU time spent (ms)=6120
  Physical memory (bytes) snapshot=795328512
  Virtual memory (bytes) snapshot=5778935808
  Total committed heap usage (bytes)=523763712

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDCOE=0
  File Input Format Counters
    Bytes Read=10000000
  File Output Format Counters
    Bytes Written=10000000
20/09/12 08:32:30 INFO TeraSort.TeraSort: done
[jessechan@instance-1 ~]$ hadoop15 bin/hadoop jar share/hadoop/mapreduce/hadoop-mapreduce-examples-2.9.2.jar teravalidate terasort/output terasort/check
20/09/12 08:32:50 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
20/09/12 08:32:51 INFO input.FileInputFormat: Total input files to process : 1
Spent 17ms computing base-splits.
Spent 3ms computing TeraScheduler splits.
20/09/12 08:32:51 INFO mapreduce.JobSubmitter: number of splits:1
20/09/12 08:32:51 INFO Configuration.deprecation: yarn.resourcemanager.system-metrics-publisher.enabled is deprecated. Instead, use yarn.system-metrics-publisher.enable
d
20/09/12 08:32:51 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1599898159880_0007
20/09/12 08:32:51 INFO impl.YarnClientImpl: Submitted application application_1599898159880_0007
20/09/12 08:32:52 INFO mapreduce.Job: The url to track the job: http://instance-1:8080/proxy/application_1599898159880_0007/
20/09/12 08:32:52 INFO mapreduce.Job: Running job: job_1599898159880_0007
20/09/12 08:32:59 INFO mapreduce.Job: Job job_1599898159880_0007 running in uber mode : false
20/09/12 08:32:59 INFO mapreduce.Job: map 0% reduce 0%
20/09/12 08:33:05 INFO mapreduce.Job: map 100% reduce 0%
20/09/12 08:33:10 INFO mapreduce.Job: map 100% reduce 100%
20/09/12 08:33:10 INFO mapreduce.Job: Job job_1599898159880_0007 completed successfully
20/09/12 08:33:10 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=92
    FILE: Number of bytes written=397313
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=10000130
    HDFS: Number of bytes written=22
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=3519
    Total time spent by all reduces in occupied slots (ms)=3150
    Total time spent by all map tasks (ms)=3519
    Total time spent by all reduce tasks (ms)=3150
    Total vcore-milliseconds taken by all map tasks=3519
    Total megabyte-milliseconds taken by all map tasks=3603456
    Total megabyte-milliseconds taken by all reduce tasks=3225600
  Map-Reduce Framework
    Map input records=100000
    Map output records=3
    Map output bytes=90
    Map output materialized bytes=92
    Input split bytes=130
    Combine input records=0
    Combine output records=0
    Reduce input groups=3
    Bytes Written=22
[jessechan@instance-1 ~]$ hadoop15 bin/hadoop dfs -ls terasort/check
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.

Found 2 items
-rw-r--r-- 3 jessechan supergroup 0 2020-09-12 08:33 terasort/check/_SUCCESS
-rw-r--r-- 3 jessechan supergroup 22 2020-09-12 08:33 terasort/check/part-r-00000
```

Use command to validate the sorting result,

```
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDCOE=0
File Input Format Counters
  Bytes Read=10000000
File Output Format Counters
  Bytes Written=22
[jessechan@instance-1 ~]$ hadoop15 bin/hadoop dfs -ls terasort/check
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.

Found 2 items
-rw-r--r-- 3 jessechan supergroup 0 2020-09-12 08:33 terasort/check/_SUCCESS
-rw-r--r-- 3 jessechan supergroup 22 2020-09-12 08:33 terasort/check/part-r-00000
```


Multi-node Hadoop Cluster Setup

Using One name-node as master and 3 data-node as slave

Datanode Information

✓ In service ⬇ Down ⚠ Decommissioned ⚠ Decommissioned & dead 🛠 In Maintenance & dead

Datanode usage histogram



In operation

Show 25 entries


Node	Http Address	Last contact	Last Block Report	Capacity	Blocks	Block pool used	Version
datanode1-1.asia-east2-a.c.driven-plexus-289109.internal.50010 (10.176.0.15:50010)	http://datanode1-1.asia-east2-a.c.driven-plexus-289109.internal:50075	2s	1m	95.75 GB	0	24 KB (0%)	2.9.2
datanode2-1.asia-east2-a.c.driven-plexus-289109.internal.50010 (10.176.0.13:50010)	http://datanode2-1.asia-east2-a.c.driven-plexus-289109.internal:50075	2s	1m	95.75 GB	0	24 KB (0%)	2.9.2
datanode3-1.asia-east2-a.c.driven-plexus-289109.internal.50010 (10.176.0.14:50010)	http://datanode3-1.asia-east2-a.c.driven-plexus-289109.internal:50075	2s	0m	95.75 GB	0	24 KB (0%)	2.9.2
namenode1.asia-east2-a.c.driven-plexus-289109.internal.50010 (10.176.0.5:50010)	http://namenode1.asia-east2-a.c.driven-plexus-289109.internal:50075	1s	1m	95.75 GB	0	24 KB (0%)	2.9.2

Showing 1 to 4 of 4 entries

Previous 1 Next

Entering Maintenance

In the 2GB TeraSort example:



Application application_1600064201081_0001

Logged in as: dr:who

Cluster

About

Nodes

Node Labels

Applications

NEW

NEW SAVING

SUBMITTED

ACCEPTED

RUNNING

FINISHED

FAILED

KILLED

Scheduler

Tools

Application Overview

User: hduiser

Name: TeraGen

Application Type: MAPREDUCE

Application Tags:

Application Priority: 0 (Higher Integer value indicates higher priority)

YarnApplicationState: FINISHED

Queue: default

FinalStatus Reported by AM: SUCCEEDED

Started: Mon Sep 14 06:17:02 +0000 2020

Elapsed: 35sec

Tracking URL: History

Log Aggregation Status: DISABLED

Application Timeout (Remaining Time): Unlimited

Diagnostics:

Unmanaged Application: false

Application Node Label expression: <Not set>

AM container Node Label expression: <DEFAULT_PARTITION>

Application Metrics

Total Resource Preempted: <memory 0, vCores 0>

Total Number of Non-AM Containers Preempted: 0

Total Number of AM Containers Preempted: 0

Resource Preempted from Current Attempt: <memory 0, vCores 0>

Number of Non-AM Containers Preempted from Current Attempt: 0

Aggregate Resource Allocation: 138496 MB-seconds, 93 vcore-seconds

Aggregate Preempted Resource Allocation: 0 MB-seconds, 0 vcore-seconds

Show 20 entries

Attempt ID

Started

Node

Logs

Nodes blacklisted by the app

Nodes blacklisted by the system

Search:

appattempt_1600064201081_0001_000001	Mon Sep 14 14:17:02 +0800 2020	http://datanode1-1.asia-east2-a.c.driven-plexus-289109.internal:8042	Logs	0	0	
--------------------------------------	--------------------------------	--	------	---	---	--

Showing 1 to 1 of 1 entries

First Previous 1 Next Last

Figure 1 2GB Gen.

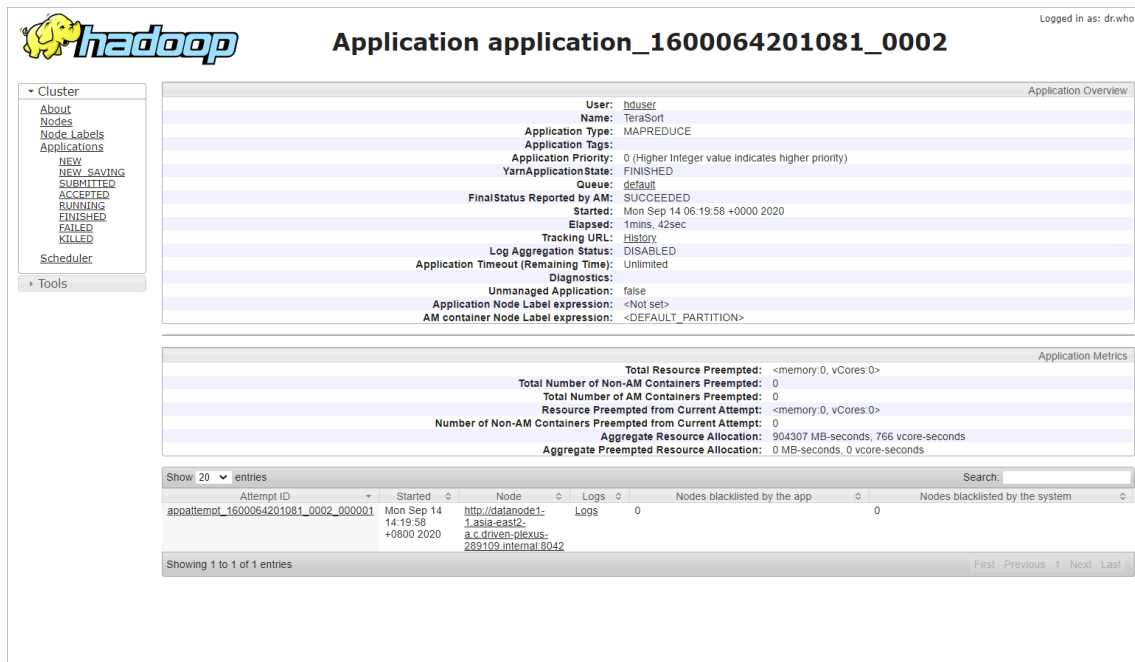


Figure 2 2GB Sort

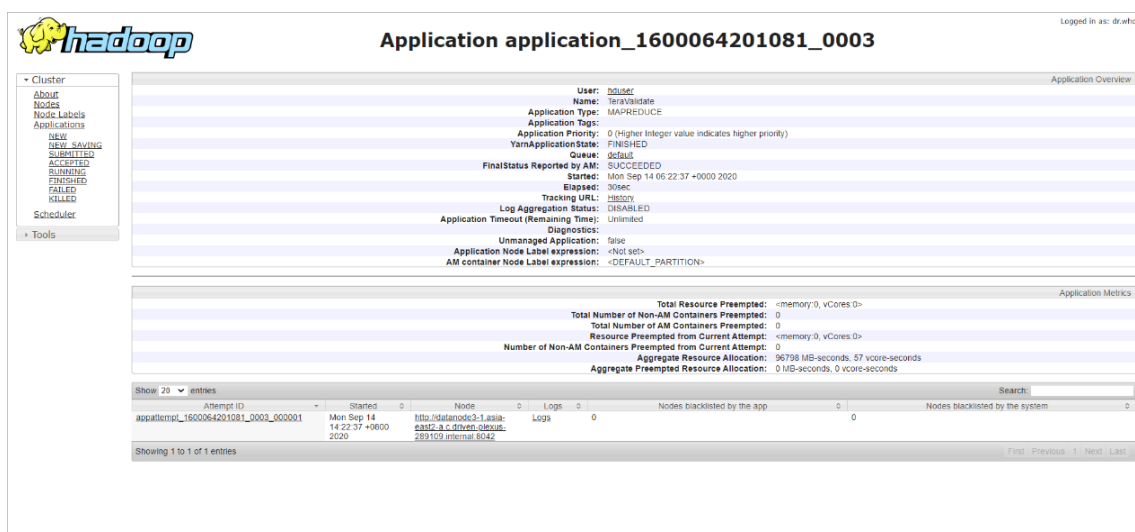


Figure 3 2GB Check

In the 20GB TeraSort example:

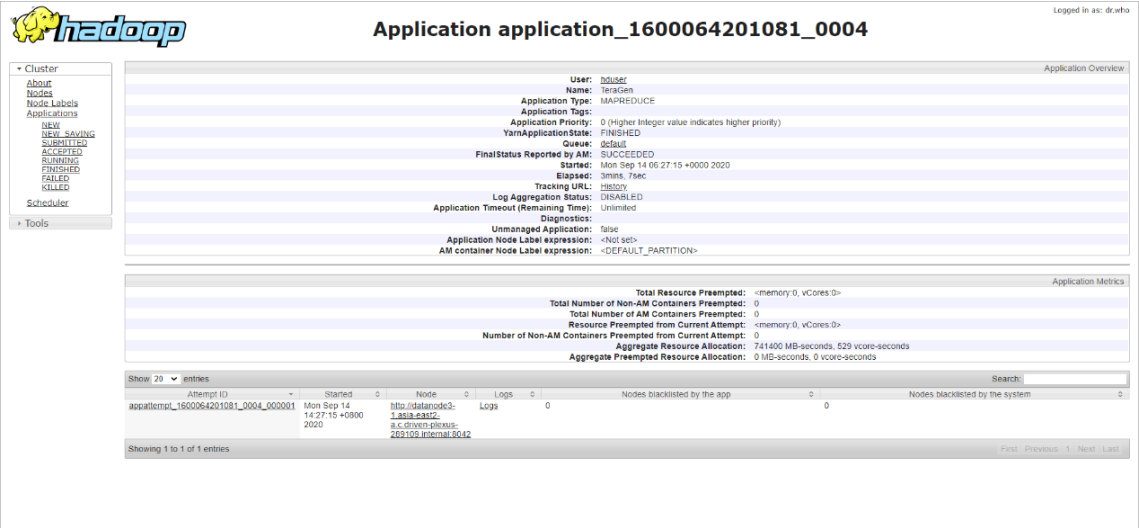


Figure 4 20GB Gen.

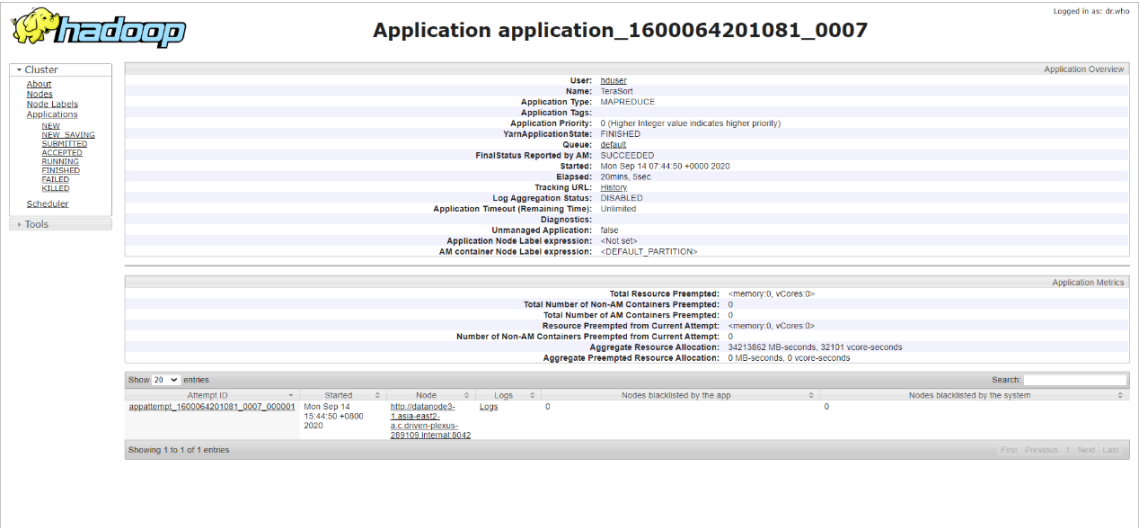


Figure 5 20GB Sort.


```
hadoop@hadoop:~$ cat /etc/hadoop/conf/hdfs-site.xml
<?xml version='1.0'?>
<configuration>
<property>
<name>dfs.replication</name>
<value>3</value>
</property>
<property>
<name>dfs.namenode.name.dir</name>
<value>/usr/local/hadoop/dfs/name</value>
</property>
<property>
<name>dfs.datanode.data.dir</name>
<value>/usr/local/hadoop/dfs/data</value>
</property>
<property>
<name>dfs.namenode.secondary.name.dir</name>
<value>/usr/local/hadoop/dfs/secondary</value>
</property>
<property>
<name>dfs.replication</name>
<value>3</value>
</property>
<property>
<name>dfs.namenode.name.dir</name>
<value>/usr/local/hadoop/dfs/name</value>
</property>
<property>
<name>dfs.datanode.data.dir</name>
<value>/usr/local/hadoop/dfs/data</value>
</property>
<property>
<name>dfs.namenode.secondary.name.dir</name>
<value>/usr/local/hadoop/dfs/secondary</value>
</property>
</configuration>
</pre>
```

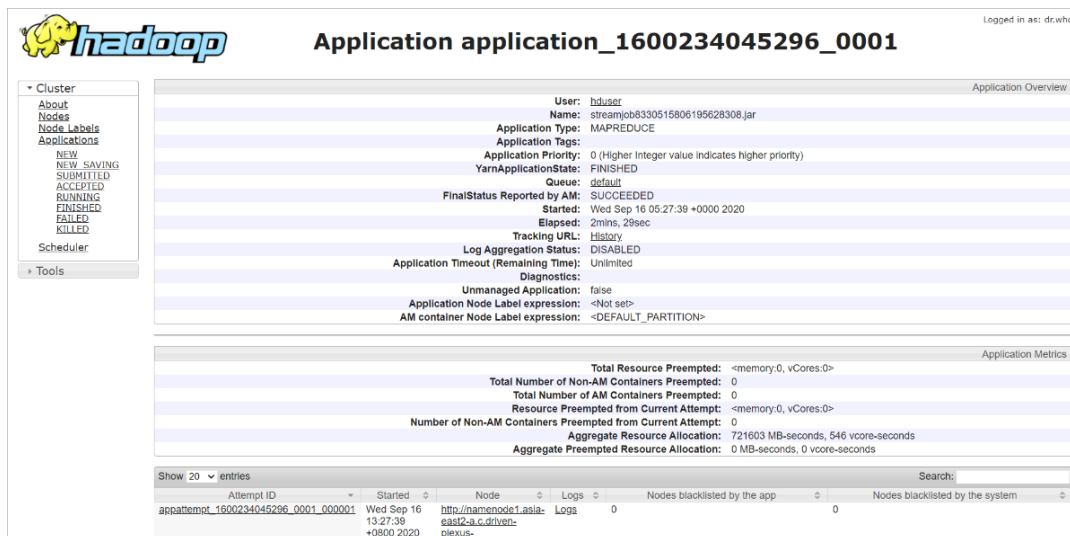


Figure 7 Python Wordcount Program

Compiling the Java WordCount program for MapReduce

```
hadoop@hadoop:~$ cat /etc/hadoop/conf/hdfs-site.xml
<?xml version='1.0'?>
<configuration>
<property>
<name>dfs.replication</name>
<value>3</value>
</property>
<property>
<name>dfs.namenode.name.dir</name>
<value>/usr/local/hadoop/dfs/name</value>
</property>
<property>
<name>dfs.datanode.data.dir</name>
<value>/usr/local/hadoop/dfs/data</value>
</property>
<property>
<name>dfs.namenode.secondary.name.dir</name>
<value>/usr/local/hadoop/dfs/secondary</value>
</property>
</configuration>
</pre>
```



```
hduser@namenode1: ~ - Google Chrome
ssh.cloud.google.com/projects/driven-plexus-289109/zones/asia-east2-a/instances/namenode1?useAdminProxy=true&authuser=0&hl=en_GB&projectNumber=743779804069

zulfikar 11
zulu 32
zululand 22
zulus 2
zum 13
zumbo 13
zumt 3
zumteln 3
zumteinspitze 3
zungeru 3
zuni 4
zuntz 4
zunyan 8
zupitza 14
zupo 2
zur 111
zurbriggen 2
zurich 77
zusammenhang 4
zutphen 8
zwaarteberg 9
zwenger'd 43
zwei 2
zwanzig 8
zweckmassigkeit 1
zwei 3
zwickau 16
zwingli 1
zwinglians 3
zwinglius 4
zischbergen 1
zwischen 1
zwolf 7
zwolferkofel 1
zwolle 4
zygema 4
zygnemaceae 7
zygogonium 5
zygoma 3
zygospore 31
zygospores 1
zygote 10
zyryanovsk 1
hduser@namenode1:~$ 2
```

The result is identical to the above python program.

Cluster

About

Nodes

Node Labels

Applications

NEW

NEW SAVING

SUBMITTED

ACCEPTED

RUNNING

FINISHED

FAILED

KILLED

Scheduler

Tools

Application Overview

User: hduser

Name: word count

Application Type: MAPREDUCE

Application Tags:

Application Priority: 0 (Higher Integer value indicates higher priority)

YarnApplicationState: FINISHED

Queue: default

FinalStatus Reported by AM: SUCCEEDED

Started: Wed Sep 16 06:40:36 +0000 2020

Elapsed: 1mins, 6sec

Tracking URL: History

Log Aggregation Status: DISABLED

Application Timeout (Remaining Time): Unlimited

Diagnostics:

Unmanaged Application: false

Application Node Label expression: <Not set>

AM container Node Label expression: <DEFAULT_PARTITION>

Application Metrics

Total Resource Preempted: <memory:0, vCores:0>

Total Number of Non-AM Containers Preempted: 0

Total Number of AM Containers Preempted: 0

Resource Preempted from Current Attempt: <memory:0, vCores:0>

Number of Non-AM Containers Preempted from Current Attempt: 0

Aggregate Resource Allocation: 263427 MB-seconds, 183 vcore-seconds

Aggregate Preempted Resource Allocation: 0 MB-seconds, 0 vcore-seconds

Show 20 entries

Attempt ID

Started

Node

Logs

Nodes blacklisted by the app

Nodes blacklisted by the system

Search:

appattemp1_1600234045296_0002_000001	Wed Sep 16 14:40:36 +0800 2020	http://namenode1.asia-east2-a.c.driven-plexus-289109.internal:8042	Logs	0	0	
--------------------------------------	--------------------------------	--	------	---	---	--

Showing 1 to 1 of 1 entries

First Previous 1 Next Last

Figure 8 Java Wordcount Program

As a conclusion, the java version of Wordcount program run faster than python program. I guess it due to Java is less dynamic than Python and It makes it more efficient on VM.